

```
!pip install pmdarima -q
!pip install tensorflow
```

```
Requirement already satisfied: tensorflow in c:\users\rishav raj
singh\anaconda3\lib\site-packages (2.16.1)
Requirement already satisfied: tensorflow-intel==2.16.1 in c:\users\
rishav raj singh\anaconda3\lib\site-packages (from tensorflow)
(2.16.1)
Requirement already satisfied: absl-py>=1.0.0 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (2.1.0)
Requirement already satisfied: astunparse>=1.6.0 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=23.5.26 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (24.3.25)
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in
c:\users\rishav raj singh\anaconda3\lib\site-packages (from
tensorflow-intel==2.16.1->tensorflow) (0.5.4)
Requirement already satisfied: google-pasta>=0.1.1 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (0.2.0)
Requirement already satisfied: h5py>=3.10.0 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (3.10.0)
Requirement already satisfied: libclang>=13.0.0 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (18.1.1)
Requirement already satisfied: ml-dtypes~=0.3.1 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (0.3.2)
Requirement already satisfied: opt-einsum>=2.3.2 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (3.3.0)
Requirement already satisfied: packaging in c:\users\rishav raj singh\
anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (23.1)
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!
=4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (3.20.3)
Requirement already satisfied: requests<3,>=2.21.0 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (2.31.0)
Requirement already satisfied: setuptools in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (68.2.2)
Requirement already satisfied: six>=1.12.0 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
```

```
>tensorflow) (1.16.0)
Requirement already satisfied: termcolor>=1.1.0 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (2.4.0)
Requirement already satisfied: typing-extensions>=3.6.6 in c:\users\
rishav raj singh\anaconda3\lib\site-packages (from tensorflow-
intel==2.16.1->tensorflow) (4.10.0)
Requirement already satisfied: wrapt>=1.11.0 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (1.14.1)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (1.62.1)
Requirement already satisfied: tensorboard<2.17,>=2.16 in c:\users\
rishav raj singh\anaconda3\lib\site-packages (from tensorflow-
intel==2.16.1->tensorflow) (2.16.2)
Requirement already satisfied: keras>=3.0.0 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (3.1.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in
c:\users\rishav raj singh\anaconda3\lib\site-packages (from
tensorflow-intel==2.16.1->tensorflow) (0.31.0)
Requirement already satisfied: numpy<2.0.0,>=1.23.5 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1-
>tensorflow) (1.26.4)
Requirement already satisfied: wheel<1.0,>=0.23.0 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from astunparse>=1.6.0-
>tensorflow-intel==2.16.1->tensorflow) (0.41.2)
Requirement already satisfied: rich in c:\users\rishav raj singh\
anaconda3\lib\site-packages (from keras>=3.0.0->tensorflow-
intel==2.16.1->tensorflow) (13.3.5)
Requirement already satisfied: namex in c:\users\rishav raj singh\
anaconda3\lib\site-packages (from keras>=3.0.0->tensorflow-
intel==2.16.1->tensorflow) (0.0.7)
Requirement already satisfied: optree in c:\users\rishav raj singh\
anaconda3\lib\site-packages (from keras>=3.0.0->tensorflow-
intel==2.16.1->tensorflow) (0.11.0)
Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\
rishav raj singh\anaconda3\lib\site-packages (from
requests<3,>=2.21.0->tensorflow-intel==2.16.1->tensorflow) (2.0.4)
Requirement already satisfied: idna<4,>=2.5 in c:\users\rishav raj
singh\anaconda3\lib\site-packages (from requests<3,>=2.21.0-
>tensorflow-intel==2.16.1->tensorflow) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from requests<3,>=2.21.0-
>tensorflow-intel==2.16.1->tensorflow) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\rishav
raj singh\anaconda3\lib\site-packages (from requests<3,>=2.21.0-
>tensorflow-intel==2.16.1->tensorflow) (2024.2.2)
```

Requirement already satisfied: markdown>=2.6.8 in c:\users\rishav raj singh\anaconda3\lib\site-packages (from tensorboard<2.17,>=2.16->tensorflow-intel==2.16.1->tensorflow) (3.4.1)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in c:\users\rishav raj singh\anaconda3\lib\site-packages (from tensorboard<2.17,>=2.16->tensorflow-intel==2.16.1->tensorflow) (0.7.2)

Requirement already satisfied: werkzeug>=1.0.1 in c:\users\rishav raj singh\anaconda3\lib\site-packages (from tensorboard<2.17,>=2.16->tensorflow-intel==2.16.1->tensorflow) (2.2.3)

Requirement already satisfied: MarkupSafe>=2.1.1 in c:\users\rishav raj singh\anaconda3\lib\site-packages (from werkzeug>=1.0.1->tensorboard<2.17,>=2.16->tensorflow-intel==2.16.1->tensorflow) (2.1.3)

Requirement already satisfied: markdown-it-py<3.0.0,>=2.2.0 in c:\users\rishav raj singh\anaconda3\lib\site-packages (from rich->keras>=3.0.0->tensorflow-intel==2.16.1->tensorflow) (2.2.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in c:\users\rishav raj singh\anaconda3\lib\site-packages (from rich->keras>=3.0.0->tensorflow-intel==2.16.1->tensorflow) (2.15.1)

Requirement already satisfied: mdurl~=0.1 in c:\users\rishav raj singh\anaconda3\lib\site-packages (from markdown-it-py<3.0.0,>=2.2.0->rich->keras>=3.0.0->tensorflow-intel==2.16.1->tensorflow) (0.1.0)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from datetime import datetime
import seaborn as sns
import plotly.express as px
from itertools import product
import warnings
import statsmodels.api as sm
import statsmodels.api as sm
sns.set_style("darkgrid")
```

```
BTC_DF = pd.read_csv("bitstampUSD_1-min_data_2012-01-01_to_2021-03-31.csv")
```

```
BTC_DF.head()
```

	Timestamp	Open	High	Low	Close	Volume_(BTC)
0	1325317920	4.39	4.39	4.39	4.39	0.455581
1	1325317980	NaN	NaN	NaN	NaN	NaN
2	1325318040	NaN	NaN	NaN	NaN	NaN
3	1325318100	NaN	NaN	NaN	NaN	NaN

```
4 1325318160    NaN    NaN    NaN    NaN    NaN
NaN
```

```
   Weighted_Price
0             4.39
1              NaN
2              NaN
3              NaN
4              NaN
```

```
BTC_DF.head()
```

```
   Timestamp  Open  High  Low  Close  Volume_(BTC)
Volume_(Currency) \
0 1325317920  4.39  4.39  4.39  4.39    0.455581
2.0
1 1325317980   NaN   NaN   NaN   NaN         NaN
NaN
2 1325318040   NaN   NaN   NaN   NaN         NaN
NaN
3 1325318100   NaN   NaN   NaN   NaN         NaN
NaN
4 1325318160   NaN   NaN   NaN   NaN         NaN
NaN
```

```
   Weighted_Price
0             4.39
1              NaN
2              NaN
3              NaN
4              NaN
```

```
# Converting the Timestamp column from string to datetime
BTC_DF['Timestamp'] = [datetime.fromtimestamp(x) for x in
BTC_DF['Timestamp']]
```

```
BTC_DF.shape
```

```
(4857377, 8)
```

```
BTC_DF.set_index("Timestamp").Weighted_Price.plot(figsize=(14,7),
title="Bitcoin Weighted Price")
```

```
<Axes: title={'center': 'Bitcoin Weighted Price'}, xlabel='Timestamp'>
```



*#calculating missing values in the dataset*

```
missing_values = BTC_DF.isnull().sum()
missing_per = (missing_values/BTC_DF.shape[0])*100
missing_table = pd.concat([missing_values,missing_per], axis=1,
ignore_index=True)
missing_table.rename(columns={0:'Total Missing Values',1:'Missing %'},
inplace=True)
missing_table
```

	Total Missing Values	Missing %
Timestamp	0	0.00000
Open	1243608	25.60246
High	1243608	25.60246
Low	1243608	25.60246
Close	1243608	25.60246
Volume_(BTC)	1243608	25.60246
Volume_(Currency)	1243608	25.60246
Weighted_Price	1243608	25.60246

```
def fill_missing(df):
    ### function to impute missing values using interpolation ###
    df['Open'] = df['Open'].interpolate()
    df['Close'] = df['Close'].interpolate()
    df['Weighted_Price'] = df['Weighted_Price'].interpolate()

    df['Volume_(BTC)'] = df['Volume_(BTC)'].interpolate()
    df['Volume_(Currency)'] = df['Volume_(Currency)'].interpolate()
    df['High'] = df['High'].interpolate()
    df['Low'] = df['Low'].interpolate()
```

```
print(df.head())
print(df.isnull().sum())
```

```
fill_missing(BTC_DF)
```

	Timestamp	Open	High	Low	Close	Volume_(BTC)	\
0	2011-12-31 13:22:00	4.39	4.39	4.39	4.39	0.455581	
1	2011-12-31 13:23:00	4.39	4.39	4.39	4.39	0.555046	
2	2011-12-31 13:24:00	4.39	4.39	4.39	4.39	0.654511	
3	2011-12-31 13:25:00	4.39	4.39	4.39	4.39	0.753977	
4	2011-12-31 13:26:00	4.39	4.39	4.39	4.39	0.853442	

	Volume_(Currency)	Weighted_Price
0	2.000000	4.39
1	2.436653	4.39
2	2.873305	4.39
3	3.309958	4.39
4	3.746611	4.39

```
Timestamp      0
Open           0
High           0
Low            0
Close          0
Volume_(BTC)   0
Volume_(Currency) 0
Weighted_Price 0
dtype: int64
```

```
#created a copy
```

```
bitstamp_non_indexed = BTC_DF.copy()
```

```
BTC_DF = BTC_DF.set_index('Timestamp')
```

```
BTC_DF.head()
```

	Open	High	Low	Close	Volume_(BTC)
Volume_(Currency) \					
Timestamp					
2011-12-31 13:22:00	4.39	4.39	4.39	4.39	0.455581
2.000000					
2011-12-31 13:23:00	4.39	4.39	4.39	4.39	0.555046
2.436653					
2011-12-31 13:24:00	4.39	4.39	4.39	4.39	0.654511
2.873305					
2011-12-31 13:25:00	4.39	4.39	4.39	4.39	0.753977
3.309958					
2011-12-31 13:26:00	4.39	4.39	4.39	4.39	0.853442
3.746611					


	Weighted_Price
Timestamp	

2011-12-31	13:22:00	4.39
2011-12-31	13:23:00	4.39
2011-12-31	13:24:00	4.39
2011-12-31	13:25:00	4.39
2011-12-31	13:26:00	4.39

```
ax = BTC_DF['Weighted_Price'].plot(title='Bitcoin Prices', grid=True,
figsize=(14,7))
ax.set_xlabel('Year')
ax.set_ylabel('Weighted Price')
```

```
ax.axvspan('2018-12-01','2019-01-31',color='red', alpha=0.3)
ax.axhspan(17500,20000, color='green',alpha=0.3)
```


<matplotlib.patches.Polygon at 0x2e38d6dd510>



*#Zooming in*

```
ax = BTC_DF.loc['2017-10':'2019-03','Weighted_Price'].plot(marker='o',
linestyle='-',figsize=(15,6), title="Oct-17 to March-19 Trend",
grid=True)
ax.set_xlabel('Month')
ax.set_ylabel('Weighted_Price')
```

```
Text(0, 0.5, 'Weighted_Price')
```



```
sns.kdeplot(BTC_DF['Weighted_Price'], shade=True)
```

```
C:\Users\RISHAV RAJ SINGH\AppData\Local\Temp\ipykernel_44116\1539549472.py:1: FutureWarning:
```


```
`shade` is now deprecated in favor of `fill`; setting `fill=True`.  
This will become an error in seaborn v0.14.0; please update your code.
```

```
sns.kdeplot(BTC_DF['Weighted_Price'], shade=True)  
C:\Users\RISHAV RAJ SINGH\anaconda3\Lib\site-packages\seaborn\oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated  
and will be removed in a future version. Convert inf values to NaN  
before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```

```
<Axes: xlabel='Weighted_Price', ylabel='Density'>
```





```
plt.figure(figsize=(15,12))
plt.suptitle('Lag Plots', fontsize=22)

plt.subplot(3,3,1)
pd.plotting.lag_plot(BTC_DF['Weighted_Price'], lag=1) #minute lag
plt.title('1-Minute Lag')

plt.subplot(3,3,2)
pd.plotting.lag_plot(BTC_DF['Weighted_Price'], lag=60) #hourley lag
plt.title('1-Hour Lag')

plt.subplot(3,3,3)
pd.plotting.lag_plot(BTC_DF['Weighted_Price'], lag=1440) #Daily lag
plt.title('Daily Lag')

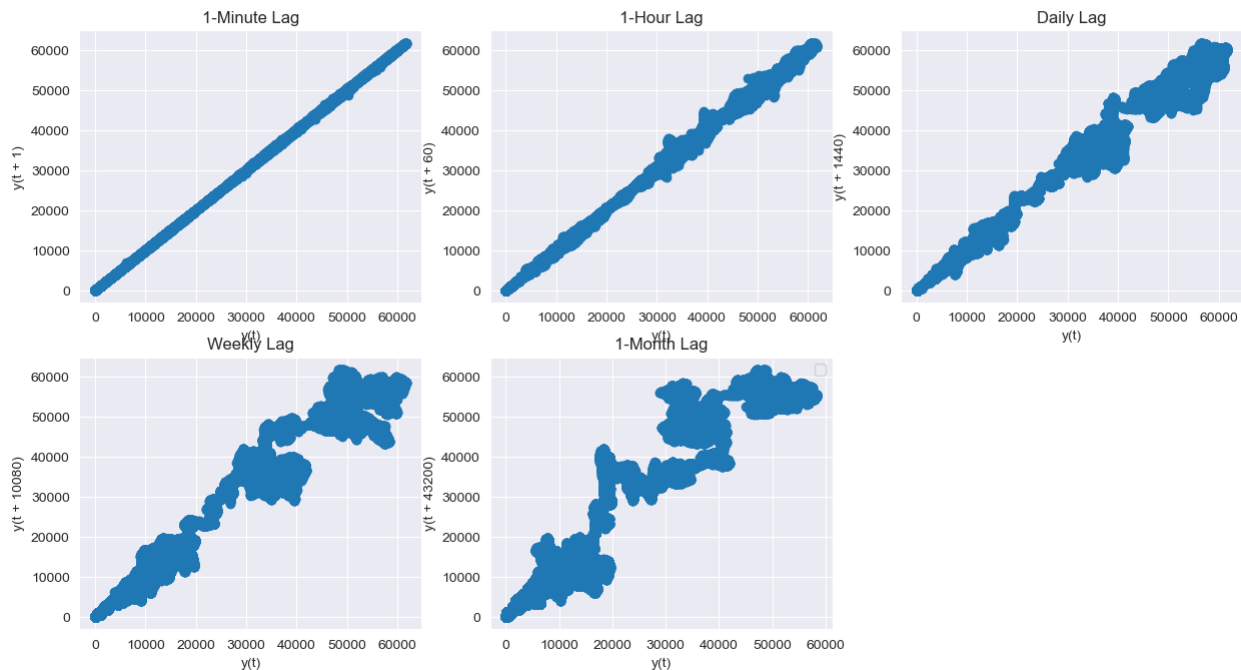
plt.subplot(3,3,4)
pd.plotting.lag_plot(BTC_DF['Weighted_Price'], lag=10080) #weekly lag
plt.title('Weekly Lag')

plt.subplot(3,3,5)
pd.plotting.lag_plot(BTC_DF['Weighted_Price'], lag=43200) #month lag
plt.title('1-Month Lag')

plt.legend()
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

## Lag Plots



```
hourly_data = BTC_DF.resample('1H').mean()
hourly_data = hourly_data.reset_index()
```

```
hourly_data.head()
```

	Timestamp	Open	High	Low	Close	Volume_(BTC)	\
0	2011-12-31 13:00:00	4.39	4.39	4.39	4.39	2.295689	
1	2011-12-31 14:00:00	4.39	4.39	4.39	4.39	7.169489	
2	2011-12-31 15:00:00	4.39	4.39	4.39	4.39	13.137408	
3	2011-12-31 16:00:00	4.39	4.39	4.39	4.39	19.105327	
4	2011-12-31 17:00:00	4.39	4.39	4.39	4.39	25.073246	

	Volume_(Currency)	Weighted_Price
0	10.078075	4.39
1	31.474059	4.39
2	57.673222	4.39
3	83.872385	4.39
4	110.071548	4.39

```
BTC_Price_daily = BTC_DF.resample("24H").mean() #daily resampling
```


```

import plotly.express as px

BTC_Price_daily.reset_index(inplace=True)
fig = px.line(BTC_Price_daily, x='Timestamp', y='Weighted_Price',
title='Weighted Price with Range Slider and Selectors')
fig.update_layout(hovermode="x")

fig.update_xaxes(
    rangeslider_visible=True,
    rangeselector=dict(
        buttons=list([
            dict(count=1, label="1m", step="month",
stepmode="backward"),
            dict(count=6, label="6m", step="month",
stepmode="backward"),
            dict(count=1, label="1y", step="year",
stepmode="backward"),
            dict(count=2, label="2y", step="year",
stepmode="backward"),
            dict(step="all")
        ])
    )
)
fig.show()

```



### #TIME SERIES DECOMPOSITION AND STATICAL TESTS

```

from statsmodels.tsa.seasonal import seasonal_decompose

from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

fill_missing(BTC_Price_daily)

```

	Timestamp	Open	High	Low	Close	Volume_(BTC)
0	2011-12-31	4.422837	4.426677	4.422837	4.426677	24.024874

1	2012-01-01	4.677625	4.677625	4.677625	4.677625	5.883361
2	2012-01-02	4.991701	4.991701	4.991701	4.991701	13.503075
3	2012-01-03	5.175495	5.175495	5.175495	5.175495	11.136690
4	2012-01-04	5.120500	5.133742	5.120500	5.133742	8.327158

	Volume_(Currency)	Weighted_Price
0	105.980529	4.424286
1	27.923145	4.677625
2	67.432386	4.991701
3	56.749845	5.175495
4	43.510443	5.125202

Timestamp	0
Open	0
High	0
Low	0
Close	0
Volume_(BTC)	0
Volume_(Currency)	0
Weighted_Price	0

dtype: int64

BTC\_Price\_daily.Weighted\_Price

0	4.424286
1	4.677625
2	4.991701
3	5.175495
4	5.125202

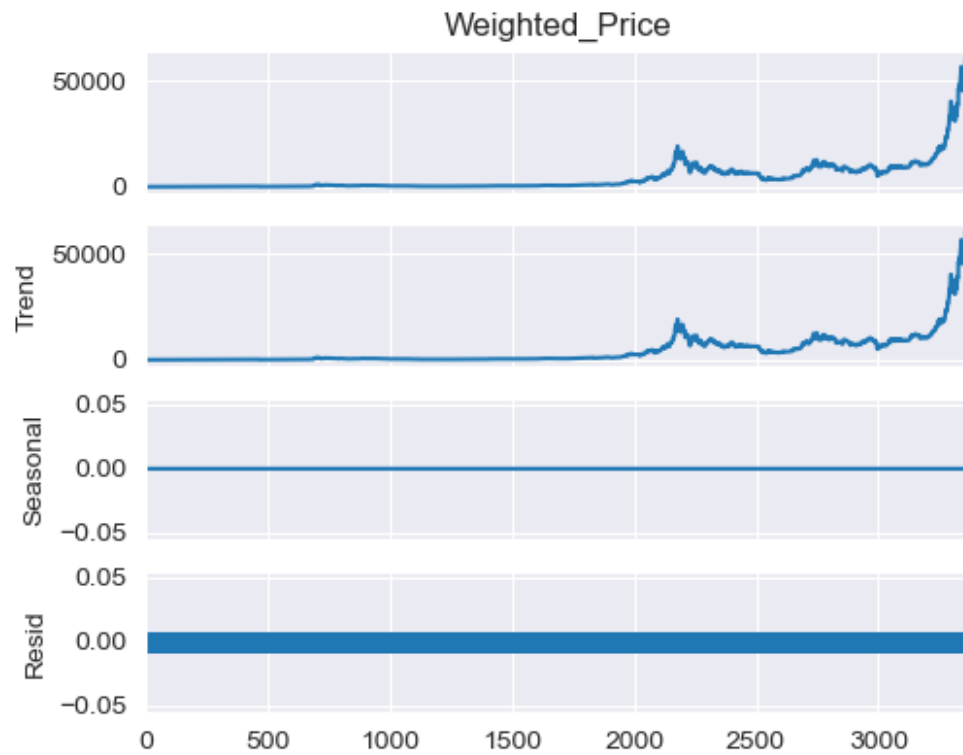
...

3374	54788.954020
3375	56002.734323
3376	56376.937694
3377	58075.416823
3378	58758.891360


Name: Weighted\_Price, Length: 3379, dtype: float64

```
plt.figure(figsize=(15,12))
series = BTC_Price_daily.Weighted_Price
result = seasonal_decompose(series, model='additive',period=1)
result.plot().set_size_inches(5, 4)
```


<Figure size 1500x1200 with 0 Axes>



```
acf = plot_acf(series, lags=50, alpha=0.05)
plt.title("ACF for Weighted Price", size=20,)
plt.show()
```



```
plot_pacf(series, lags=50, alpha=0.05, method='ols')  
plt.title("PACF for Weighted Price", size=20)  
plt.show()
```



```
#stationary time series
```

```
from statsmodels.tsa.stattools import kpss
```

```
#adf test
```

```
stats, p, lags, critical_values = kpss(series, 'ct')
```

```
C:\Users\RISHAV RAJ SINGH\AppData\Local\Temp\
ipykernel_44116\108484339.py:2: InterpolationWarning:
```

The test statistic is outside of the range of p-values available in the look-up table. The actual p-value is smaller than the p-value returned.

```
def adf_test(timeseries):
    print ('Results of Dickey-Fuller Test:')
    dfctest = adfuller(timeseries, autolag='AIC')
    print("dfctest: ", dfctest)
    dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic','p-
value','#Lags Used','Number of Observations Used'])
    for key,value in dfctest[4].items():
        dfoutput['Critical Value (%s)'%key] = value
```

```

print (dfoutput)

if p > 0.05:
    print('Series is not Stationary')
else:
    print('Series is Stationary')

adf_test(series)

```

Results of Dickey-Fuller Test:

```

dfctest: (6.302120534312761, 1.0, 29, 3349, {'1%': -3.432304111473485,
'5%': -2.862403412310526, '10%': -2.56722961145352},
48472.31981125205)
Test Statistic                6.302121
p-value                      1.000000
#Lags Used                    29.000000
Number of Observations Used   3349.000000
Critical Value (1%)           -3.432304
Critical Value (5%)           -2.862403
Critical Value (10%)          -2.567230
dtype: float64
Series is Stationary

```

### #FEATURE ENGINEERING

```

df = BTC_Price_daily.set_index("Timestamp")
df.reset_index(drop=False, inplace=True)

rolling_features = ["Open", "High", "Low", "Close", "Volume_(BTC)"]
window1 = 3
window2 = 7
window3 = 30

# First convert our original df to a rolling df of 3d, 7d and 30d
df_rolled_3d = df[rolling_features].rolling(window=window1,
min_periods=0)
df_rolled_7d = df[rolling_features].rolling(window=window2,
min_periods=0)
df_rolled_30d = df[rolling_features].rolling(window=window3,
min_periods=0)

# dataframe.shift() function Shift index by desired number of periods.
It takes a scalar parameter called the period,
# which represents the number of shifts to be made over the desired
axis. It defaults to 1 and
# it is shifting values vertically along the axis 0 . NaN will be
filled for missing values introduced as a result of the shifting.
# Very helpful when dealing with time-series data.

```



*# <https://towardsdatascience.com/all-the-pandas-shift-you-should-know-for-data-analysis-791c1692b5e>*

```
df_mean_3d = df_rolled_3d.mean().shift(1).reset_index()
df_mean_7d = df_rolled_7d.mean().shift(1).reset_index()
df_mean_30d = df_rolled_30d.mean().shift(1).reset_index()
```

*# Just print to see the structure of one of them*

df\_mean\_30d

	index	Open	High	Low	Close \
0	0	NaN	NaN	NaN	NaN
1	1	4.422837	4.426677	4.422837	4.426677
2	2	4.550231	4.552151	4.550231	4.552151
3	3	4.697388	4.698668	4.697388	4.698668
4	4	4.816915	4.817875	4.816915	4.817875
...	...	...	...	...	...
3374	3374	53206.447763	53245.227654	53167.556178	53207.530781
3375	3375	53364.378192	53402.510297	53326.251128	53365.529760
3376	3376	53658.228941	53694.963607	53621.648890	53659.456192
3377	3377	53968.885871	54005.334879	53932.677167	53970.075025
3378	3378	54394.961787	54430.840483	54359.421673	54396.251845

	Volume_(BTC)
0	NaN
1	24.024874
2	14.954117
3	14.470437
4	13.637000
...	...
3374	3.852066
3375	3.754108
3376	3.594949
3377	3.623054
3378	3.581530

[3379 rows x 6 columns]

```
df_std_3d = df_rolled_3d.std().shift(1).reset_index()
df_std_7d = df_rolled_7d.std().shift(1).reset_index()
df_std_30d = df_rolled_30d.std().shift(1).reset_index()
```

*# Just print to see the structure of one of them*

df\_std\_30d

	index	Open	High	Low	Close
Volume_(BTC)					
0	0	NaN	NaN	NaN	NaN
NaN					
1	1	NaN	NaN	NaN	NaN

```

NaN
2      2      0.180163      0.177447      0.180163      0.177447
12.827987
3      3      0.284946      0.283099      0.284946      0.283099
9.109362
4      4      0.333581      0.332071      0.333581      0.332071
7.622256
...      ...      ...      ...      ...      ...
...
3374    3374    4435.974803    4434.057878    4437.954547    4436.182696
1.223055
3375    3375    4403.971836    4402.115028    4405.906128    4404.182498
1.250130
3376    3376    4269.649305    4268.898044    4270.343124    4269.817028
1.245602
3377    3377    4108.821654    4108.182451    4109.416624    4109.022442
1.232919
3378    3378    3831.773514    3831.207184    3831.997540    3831.768848
1.238370

```

```
[3379 rows x 6 columns]
```

```

for feature in rolling_features:
    df[f"{feature}_mean_lag{window1}"] = df_mean_3d[feature]
    df[f"{feature}_mean_lag{window2}"] = df_mean_7d[feature]
    df[f"{feature}_mean_lag{window3}"] = df_mean_30d[feature]

    df[f"{feature}_std_lag{window1}"] = df_std_3d[feature]
    df[f"{feature}_std_lag{window2}"] = df_std_7d[feature]
    df[f"{feature}_std_lag{window3}"] = df_std_30d[feature]

```

```
df.fillna(df.mean(), inplace=True)
```

```
df.set_index("Timestamp", drop=False, inplace=True)
df.head()
```

	Timestamp	Open	High	Low	Close
Volume_(BTC) \					
Timestamp					
2011-12-31	2011-12-31	4.422837	4.426677	4.422837	4.426677
24.024874					
2012-01-01	2012-01-01	4.677625	4.677625	4.677625	4.677625
5.883361					
2012-01-02	2012-01-02	4.991701	4.991701	4.991701	4.991701
13.503075					
2012-01-03	2012-01-03	5.175495	5.175495	5.175495	5.175495
11.136690					
2012-01-04	2012-01-04	5.120500	5.133742	5.120500	5.133742
8.327158					

Open_mean_lag7 \	Volume_(Currency)	Weighted_Price	Open_mean_lag3
Timestamp			

2011-12-31	105.980529	4.424286	4564.752758
4532.014863			
2012-01-01	27.923145	4.677625	4.422837
4.422837			
2012-01-02	67.432386	4.991701	4.550231
4.550231			
2012-01-03	56.749845	5.175495	4.697388
4.697388			
2012-01-04	43.510443	5.125202	4.948274
4.816915			

Timestamp	...	Close_mean_lag30	Close_std_lag3	Close_std_lag7 \
2011-12-31	...	4342.158190	112.850620	186.071830
2012-01-01	...	4.426677	112.850620	186.071830
2012-01-02	...	4.552151	0.177447	0.177447
2012-01-03	...	4.698668	0.283099	0.283099
2012-01-04	...	4.817875	0.251760	0.332071

Volume_(BTC)_mean_lag7 \	Close_std_lag30	Volume_(BTC)_mean_lag3
Timestamp		

2011-12-31	417.775512	9.565484
9.572699		
2012-01-01	417.775512	24.024874
24.024874		
2012-01-02	0.177447	14.954117
14.954117		
2012-01-03	0.283099	14.470437
14.470437		
2012-01-04	0.332071	10.174375
13.637000		

Timestamp	Volume_(BTC)_mean_lag30	Volume_(BTC)_std_lag3 \
2011-12-31	9.590264	3.369452
2012-01-01	24.024874	3.369452
2012-01-02	14.954117	12.827987
2012-01-03	14.470437	9.109362
2012-01-04	13.637000	3.899942

Timestamp	Volume_(BTC)_std_lag7	Volume_(BTC)_std_lag30
2011-12-31	4.136303	5.128796

2012-01-01	4.136303	5.128796
2012-01-02	12.827987	12.827987
2012-01-03	9.109362	9.109362
2012-01-04	7.622256	7.622256

[5 rows x 38 columns]

```
df["month"] = df.Timestamp.dt.month
df["week"] = df.Timestamp.dt.isocalendar().week
df["day"] = df.Timestamp.dt.day
df["day_of_week"] = df.Timestamp.dt.dayofweek
df.head()
```

	Timestamp	Open	High	Low	Close
Volume_(BTC) \	Timestamp				
2011-12-31	2011-12-31	4.422837	4.426677	4.422837	4.426677
24.024874					
2012-01-01	2012-01-01	4.677625	4.677625	4.677625	4.677625
5.883361					
2012-01-02	2012-01-02	4.991701	4.991701	4.991701	4.991701
13.503075					
2012-01-03	2012-01-03	5.175495	5.175495	5.175495	5.175495
11.136690					
2012-01-04	2012-01-04	5.120500	5.133742	5.120500	5.133742
8.327158					

	Volume_(Currency)	Weighted_Price	Open_mean_lag3
Open_mean_lag7 \			
Timestamp			
2011-12-31	105.980529	4.424286	4564.752758
4532.014863			
2012-01-01	27.923145	4.677625	4.422837
4.422837			
2012-01-02	67.432386	4.991701	4.550231
4.550231			
2012-01-03	56.749845	5.175495	4.697388
4.697388			
2012-01-04	43.510443	5.125202	4.948274
4.816915			

	...	Volume_(BTC)_mean_lag3	Volume_(BTC)_mean_lag7 \
Timestamp	...		
2011-12-31	...	9.565484	9.572699
2012-01-01	...	24.024874	24.024874
2012-01-02	...	14.954117	14.954117
2012-01-03	...	14.470437	14.470437
2012-01-04	...	10.174375	13.637000

	Volume_(BTC)_mean_lag30	Volume_(BTC)_std_lag30 \
Timestamp		
2011-12-31	9.590264	3.369452
2012-01-01	24.024874	3.369452
2012-01-02	14.954117	12.827987
2012-01-03	14.470437	9.109362
2012-01-04	13.637000	3.899942

	Volume_(BTC)_std_lag7	Volume_(BTC)_std_lag30	month	week
day \				
Timestamp				
2011-12-31	4.136303	5.128796	12	52
31				
2012-01-01	4.136303	5.128796	1	52
1				
2012-01-02	12.827987	12.827987	1	1
2				
2012-01-03	9.109362	9.109362	1	1
3				
2012-01-04	7.622256	7.622256	1	1
4				

	day_of_week
Timestamp	
2011-12-31	5
2012-01-01	6
2012-01-02	0
2012-01-03	1
2012-01-04	2

[5 rows x 42 columns]

### #MODEL BUILDING

```
df_train = df[df.Timestamp < "2020"]
df_valid = df[df.Timestamp >= "2020"]

print('train shape :', df_train.shape)
print('validation shape :', df_valid.shape)

train shape : (2923, 42)
validation shape : (456, 42)

df_train = df[df.Timestamp < "2020"]
df_valid = df[df.Timestamp >= "2020"]

print('train shape :', df_train.shape)
print('validation shape :', df_valid.shape)
```

```

train shape : (2923, 42)
validation shape : (456, 42)

#ARIMA

import pmdarima as pm

# From df.columns, I remove all the original 8 columns i.e
# 'Timestamp', 'Open', 'High', 'Low', 'Close', 'Volume_(BTC)',
# 'Volume_(Currency)', 'Weighted_Price',
# So I will be left with below newly created (engineered) columns
# Also

exogenous_features = ['Open_mean_lag3',
                      'Open_mean_lag7', 'Open_mean_lag30', 'Open_std_lag3',
                      'Open_std_lag7',
                      'Open_std_lag30', 'High_mean_lag3', 'High_mean_lag7',
                      'High_mean_lag30',
                      'High_std_lag3', 'High_std_lag7', 'High_std_lag30',
                      'Low_mean_lag3',
                      'Low_mean_lag7', 'Low_mean_lag30', 'Low_std_lag3',
                      'Low_std_lag7',
                      'Low_std_lag30', 'Close_mean_lag3', 'Close_mean_lag7',
                      'Close_mean_lag30', 'Close_std_lag3', 'Close_std_lag7',
                      'Close_std_lag30', 'Volume_(BTC)_mean_lag3',
                      'Volume_(BTC)_mean_lag7',
                      'Volume_(BTC)_mean_lag30', 'Volume_(BTC)_std_lag3',
                      'Volume_(BTC)_std_lag7', 'Volume_(BTC)_std_lag30', 'month',
                      'day', 'day_of_week']

# len(exogenous_features) # 34

model = pm.auto_arima(df_train.Weighted_Price,
                      exogenous=df_train[exogenous_features], trace=True,
                      error_action="ignore", suppress_warnings=True)

model.fit(df_train.Weighted_Price,
          exogenous=df_train[exogenous_features])

forecast = model.predict(n_periods=len(df_valid),
                        exogenous=df_valid[exogenous_features])

df_valid["Forecast_ARIMAX"] = forecast

Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=38726.548, Time=2.87 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=38854.989, Time=0.12 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=38739.796, Time=0.31 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=38737.063, Time=0.31 sec
ARIMA(0,1,0)(0,0,0)[0] : AIC=38853.501, Time=0.07 sec

```

```

ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=38728.327, Time=2.10 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=38740.389, Time=0.41 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=38718.014, Time=5.14 sec
ARIMA(3,1,1)(0,0,0)[0] intercept : AIC=38742.364, Time=0.63 sec
ARIMA(4,1,2)(0,0,0)[0] intercept : AIC=38732.559, Time=2.68 sec
ARIMA(3,1,3)(0,0,0)[0] intercept : AIC=38727.437, Time=7.64 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=38725.338, Time=5.87 sec
ARIMA(4,1,1)(0,0,0)[0] intercept : AIC=38729.099, Time=1.78 sec
ARIMA(4,1,3)(0,0,0)[0] intercept : AIC=38728.781, Time=3.70 sec
ARIMA(3,1,2)(0,0,0)[0] : AIC=38716.214, Time=2.57 sec
ARIMA(2,1,2)(0,0,0)[0] : AIC=38724.804, Time=1.35 sec
ARIMA(3,1,1)(0,0,0)[0] : AIC=38730.459, Time=1.09 sec
ARIMA(4,1,2)(0,0,0)[0] : AIC=38730.822, Time=1.04 sec
ARIMA(3,1,3)(0,0,0)[0] : AIC=38725.524, Time=3.72 sec
ARIMA(2,1,1)(0,0,0)[0] : AIC=38738.753, Time=0.12 sec
ARIMA(2,1,3)(0,0,0)[0] : AIC=38730.606, Time=1.11 sec
ARIMA(4,1,1)(0,0,0)[0] : AIC=38727.378, Time=1.18 sec
ARIMA(4,1,3)(0,0,0)[0] : AIC=38727.060, Time=2.81 sec

```

Best model: ARIMA(3,1,2)(0,0,0)[0]

Total fit time: 48.654 seconds

C:\Users\RISHAV RAJ SINGH\AppData\Local\Temp\  
ipykernel\_44116\183215675.py:8: SettingWithCopyWarning:


A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_valid[["Weighted_Price", "Forecast_ARIMAX"]].plot(figsize=(15, 5))
```

```
<Axes: xlabel='Timestamp'>
```



```
from sklearn.metrics import mean_squared_error, mean_absolute_error

print("RMSE of Auto ARIMAX:",
      np.sqrt(mean_squared_error(df_valid.Weighted_Price,
                                 df_valid.Forecast_ARIMAX)))
```


```
print("\nMAE of Auto ARIMAX:",
      mean_absolute_error(df_valid.Weighted_Price,
                           df_valid.Forecast_ARIMAX))
```

RMSE of Auto ARIMAX: 18052.527118342223

MAE of Auto ARIMAX: 10692.217013361565

```
model.plot_diagnostics().set_size_inches(10, 8) # Adjust the size as
needed
plt.show()
```





*#LSTM*

```
price_series = BTC_Price_daily.reset_index().Weighted_Price.values
price_series
```

```
array([4.42428590e+00, 4.67762540e+00, 4.99170099e+00, ...,
       5.63769377e+04, 5.80754168e+04, 5.87588914e+04])
```

```
# Lets check what was the split of earlier train and test dataset
df_train.shape # (2923, 43)
```

```
(2923, 42)
```

```
train_data, test_data = price_series[0:2923], price_series[2923:]
```

```
# Feature Scaling
```

```
# from sklearn.preprocessing import MinMaxScaler
```

```
# scaler = MinMaxScaler(feature_range = (0, 1))
```

```

# price_series_scaled = scaler.fit_transform(price_series.reshape(-1,1))

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler(feature_range = (0, 1))

train_data = scaler.fit_transform(train_data.reshape(-1,1))

test_data = scaler.transform(test_data.reshape(-1,1))

# train_data, test_data = price_series_scaled[0:2923],
# price_series_scaled[2923:]

train_data.shape, test_data.shape

((2923, 1), (456, 1))

def get_window_ds_list_for_lstm(series, time_step):
    # Here, basically for each window, all data upto the last but one data-point
    # will be appended to as X variable and then
    # the last data-point of that window will be the Y variable (i.e. target)
    dataX, dataY = [], []
    for i in range(len(series)- time_step-1):
        a = series[i : (i+time_step), 0]
        dataX.append(a)
        b = series[i + time_step, 0]
        dataY.append(b)

    return np.array(dataX), np.array(dataY)

X_train, y_train = get_window_ds_list_for_lstm(train_data,
time_step=100)
X_test, y_test = get_window_ds_list_for_lstm(test_data, time_step=100)

X_train.shape, y_train.shape, X_test.shape, y_test.shape

((2822, 100), (2822,), (355, 100), (355,))

#reshape inputs to be [samples, timesteps, features] which is required for LSTM

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], 1)

print(X_train.shape)
print(X_test.shape)

```

```
(2822, 100, 1)
(355, 100, 1)
```

```
print(y_train.shape)
print(y_test.shape)
```

```
(2822,)
(355,)
```

```
#Create Stacked LSTM Model
```

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import LSTM
from tensorflow.keras.layers import Dropout
```

```
# Initialising the LSTM
regressor = Sequential()
```

```
# Adding the first LSTM layer and some Dropout regularisation
regressor.add(LSTM(units = 50, return_sequences = True, input_shape =
(X_train.shape[1], 1)))
regressor.add(Dropout(0.2))
```

```
# Adding a second LSTM layer and some Dropout regularisation
regressor.add(LSTM(units = 50, return_sequences = True))
regressor.add(Dropout(0.2))
```

```
# Adding a third LSTM layer and some Dropout regularisation
regressor.add(LSTM(units = 50, return_sequences = True))
regressor.add(Dropout(0.2))
```

```
# Adding a fourth LSTM layer and some Dropout regularisation
regressor.add(LSTM(units = 50))
regressor.add(Dropout(0.2))
```

```
# Adding the output layer
regressor.add(Dense(units = 1))
```

```
# Compiling the RNN
regressor.compile(optimizer = 'adam', loss = 'mean_squared_error')
```

```
C:\Users\RISHAV RAJ SINGH\anaconda3\Lib\site-packages\keras\src\
layers\rnn\rnn.py:204: UserWarning:
```

```
Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the
first layer in the model instead.
```

```
regressor.summary()
```

Model: "sequential"

Layer (type) Param #	Output Shape
lstm (LSTM) 10,400	(None, 100, 50)
dropout (Dropout) 0	(None, 100, 50)
lstm_1 (LSTM) 20,200	(None, 100, 50)
dropout_1 (Dropout) 0	(None, 100, 50)
lstm_2 (LSTM) 20,200	(None, 100, 50)
dropout_2 (Dropout) 0	(None, 100, 50)
lstm_3 (LSTM) 20,200	(None, 50)
dropout_3 (Dropout) 0	(None, 50)
dense (Dense) 51	(None, 1)

Total params: 71,051 (277.54 KB)

Trainable params: 71,051 (277.54 KB)

Non-trainable params: 0 (0.00 B)

```
# Fitting the RNN to the Training set
```

```
history = regressor.fit(X_train, y_train, validation_split=0.1, epochs  
= 50, batch_size = 32, verbose=1, shuffle=False)
```

```
Epoch 1/50
```

```
80/80 _____ 21s 168ms/step - loss: 9.4814e-04 -
```

```
val_loss: 0.0318
```

```
Epoch 2/50
```

```
80/80 _____ 14s 173ms/step - loss: 0.0124 - val_loss:  
0.0654
```

```
Epoch 3/50
```

```
80/80 _____ 12s 152ms/step - loss: 0.0081 - val_loss:  
0.0102
```

```
Epoch 4/50
```

```
80/80 _____ 14s 181ms/step - loss: 0.0227 - val_loss:  
0.0148
```

```
Epoch 5/50
```

```
80/80 _____ 16s 195ms/step - loss: 0.0109 - val_loss:  
0.0035
```

```
Epoch 6/50
```

```
80/80 _____ 13s 166ms/step - loss: 0.0053 - val_loss:  
0.0089
```

```
Epoch 7/50
```

```
80/80 _____ 12s 150ms/step - loss: 0.0042 - val_loss:  
0.0103
```

```
Epoch 8/50
```

```
80/80 _____ 12s 154ms/step - loss: 0.0023 - val_loss:  
0.0065
```

```
Epoch 9/50
```

```
80/80 _____ 12s 154ms/step - loss: 0.0066 - val_loss:  
0.0023
```

```
Epoch 10/50
```

```
80/80 _____ 13s 156ms/step - loss: 0.0016 - val_loss:  
0.0021
```

```
Epoch 11/50
```

```
80/80 _____ 14s 172ms/step - loss: 0.0025 - val_loss:  
0.0026
```

```
Epoch 12/50
```

```
80/80 _____ 14s 178ms/step - loss: 0.0042 - val_loss:  
0.0023
```

```
Epoch 13/50
```

```
80/80 _____ 14s 179ms/step - loss: 0.0022 - val_loss:  
0.0025
```

```
Epoch 14/50
```

```
80/80 _____ 21s 263ms/step - loss: 0.0023 - val_loss:  
0.0022
```

```
Epoch 15/50
```

```
80/80 _____ 18s 229ms/step - loss: 0.0025 - val_loss:  
0.0020
```

```
Epoch 16/50
```


```
80/80 ————— 20s 253ms/step - loss: 0.0025 - val_loss: 0.0028
Epoch 17/50
80/80 ————— 21s 256ms/step - loss: 0.0024 - val_loss: 0.0025
Epoch 18/50
80/80 ————— 19s 242ms/step - loss: 0.0022 - val_loss: 0.0041
Epoch 19/50
80/80 ————— 19s 238ms/step - loss: 0.0024 - val_loss: 0.0041
Epoch 20/50
80/80 ————— 19s 239ms/step - loss: 0.0021 - val_loss: 0.0049
Epoch 21/50
80/80 ————— 19s 224ms/step - loss: 0.0022 - val_loss: 0.0022
Epoch 22/50
80/80 ————— 19s 238ms/step - loss: 0.0018 - val_loss: 0.0022
Epoch 23/50
80/80 ————— 18s 225ms/step - loss: 0.0020 - val_loss: 0.0011
Epoch 24/50
80/80 ————— 18s 228ms/step - loss: 0.0020 - val_loss: 0.0034
Epoch 25/50
80/80 ————— 18s 224ms/step - loss: 0.0019 - val_loss: 0.0038
Epoch 26/50
80/80 ————— 18s 227ms/step - loss: 0.0017 - val_loss: 0.0026
Epoch 27/50
80/80 ————— 18s 223ms/step - loss: 0.0016 - val_loss: 0.0036
Epoch 28/50
80/80 ————— 8s 104ms/step - loss: 0.0018 - val_loss: 0.0036
Epoch 29/50
80/80 ————— 7s 89ms/step - loss: 0.0019 - val_loss: 0.0043
Epoch 30/50
80/80 ————— 7s 88ms/step - loss: 0.0019 - val_loss: 0.0056
Epoch 31/50
80/80 ————— 7s 87ms/step - loss: 0.0019 - val_loss: 0.0055
Epoch 32/50
80/80 ————— 7s 90ms/step - loss: 0.0020 - val_loss:
```

```
0.0031
Epoch 33/50
80/80 _____ 7s 88ms/step - loss: 0.0018 - val_loss:
0.0011
Epoch 34/50
80/80 _____ 7s 86ms/step - loss: 0.0019 - val_loss:
0.0010
Epoch 35/50
80/80 _____ 7s 92ms/step - loss: 0.0023 - val_loss:
9.1630e-04
Epoch 36/50
80/80 _____ 8s 105ms/step - loss: 0.0020 - val_loss:
0.0017
Epoch 37/50
80/80 _____ 7s 92ms/step - loss: 0.0014 - val_loss:
0.0029
Epoch 38/50
80/80 _____ 8s 94ms/step - loss: 0.0013 - val_loss:
0.0055
Epoch 39/50
80/80 _____ 7s 90ms/step - loss: 0.0015 - val_loss:
0.0067
Epoch 40/50
80/80 _____ 8s 95ms/step - loss: 0.0017 - val_loss:
0.0042
Epoch 41/50
80/80 _____ 9s 115ms/step - loss: 0.0014 - val_loss:
0.0025
Epoch 42/50
80/80 _____ 9s 106ms/step - loss: 0.0017 - val_loss:
0.0018
Epoch 43/50
80/80 _____ 8s 102ms/step - loss: 0.0017 - val_loss:
0.0013
Epoch 44/50
80/80 _____ 8s 99ms/step - loss: 0.0016 - val_loss:
9.5515e-04
Epoch 45/50
80/80 _____ 8s 98ms/step - loss: 0.0027 - val_loss:
9.6075e-04
Epoch 46/50
80/80 _____ 8s 100ms/step - loss: 0.0017 - val_loss:
0.0021
Epoch 47/50
80/80 _____ 8s 97ms/step - loss: 0.0011 - val_loss:
0.0027
Epoch 48/50
80/80 _____ 8s 99ms/step - loss: 0.0013 - val_loss:
0.0034
```

```
Epoch 49/50
80/80 ━━━━━━━━━━━ 8s 97ms/step - loss: 0.0014 - val_loss: 0.0031
Epoch 50/50
80/80 ━━━━━━━━━━━ 8s 101ms/step - loss: 0.0014 - val_loss: 0.0015
```

```
plt.figure(figsize=(16,7))
plt.plot(history.history["loss"], label= "train loss")
plt.plot(history.history["val_loss"], label= "validation loss")
plt.legend()
```

<matplotlib.legend.Legend at 0x2e492f5a9d0>



```
train_predict = regressor.predict(X_train)
test_predict = regressor.predict(X_test)
```

```
89/89 ━━━━━━━━━━━ 4s 37ms/step
12/12 ━━━━━━━━━━━ 0s 25ms/step
```


```
y_train_inv = scaler.inverse_transform(y_train.reshape(-1, 1))
y_test_inv = scaler.inverse_transform(y_test.reshape(-1, 1))
```

```
train_predict_inv = scaler.inverse_transform(train_predict)
test_predict_inv = scaler.inverse_transform(test_predict)
```

```
plt.figure(figsize=(16,7))
plt.plot(y_train_inv.flatten(), marker='.', label="Actual")
plt.plot(train_predict_inv.flatten(), 'r', marker='.',
label="Predicted")
plt.legend()
```


<matplotlib.legend.Legend at 0x2e4932482d0>





```
plt.figure(figsize=(16,7))
plt.plot(y_test_inv.flatten(), marker='.', label="Actual")
plt.plot(test_predict_inv.flatten(), 'r', marker='.',
label="Predicted")
plt.legend()
```

<matplotlib.legend.Legend at 0x2e493849b10>

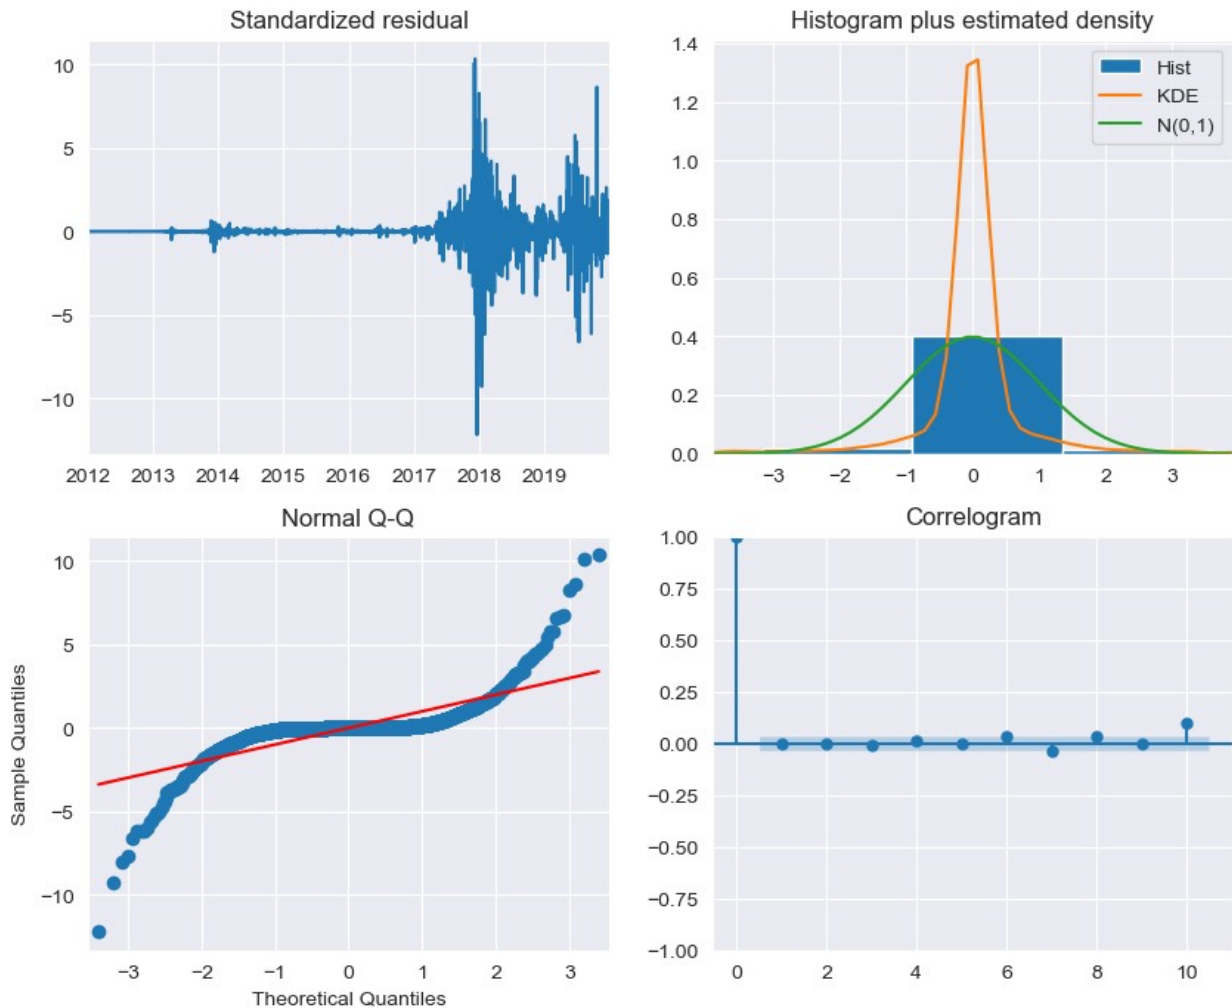


```
LSTM_train_RMSE_inverse = np.sqrt(mean_squared_error(y_train_inv,
train_predict_inv))
LSTM_test_RMSE_inverse = np.sqrt(mean_squared_error(y_test_inv,
test_predict_inv))
```

```
LSTM_train_MAE_inverse = mean_squared_error(y_train_inv,
train_predict_inv )
```

```
LSTM_test_MAE_inverse = mean_squared_error(y_test_inv,
test_predict_inv )

model.plot_diagnostics().set_size_inches(10, 8) # Adjust the size as
needed
plt.show()
```



### #PROPHET

```
from prophet import Prophet

# Resampling original data to day level and forward fill the missing
values
daily_data = BTC_DF.resample("24H").mean() #daily resampling
fill_missing(daily_data)
```

	Open	High	Low	Close	Volume_(BTC)
Timestamp					
2011-12-31	4.422837	4.426677	4.422837	4.426677	24.024874

2012-01-01	4.677625	4.677625	4.677625	4.677625	5.883361
2012-01-02	4.991701	4.991701	4.991701	4.991701	13.503075
2012-01-03	5.175495	5.175495	5.175495	5.175495	11.136690
2012-01-04	5.120500	5.133742	5.120500	5.133742	8.327158

	Volume_(Currency)	Weighted_Price
Timestamp		
2011-12-31	105.980529	4.424286
2012-01-01	27.923145	4.677625
2012-01-02	67.432386	4.991701
2012-01-03	56.749845	5.175495
2012-01-04	43.510443	5.125202
Open	0	
High	0	
Low	0	
Close	0	
Volume_(BTC)	0	
Volume_(Currency)	0	
Weighted_Price	0	

dtype: int64

*# Renaming the column names according to Prophet's requirements*

```
daily_data_fb = daily_data.reset_index()
[['Timestamp', 'Weighted_Price']].rename({'Timestamp': 'ds', 'Weighted_Price': 'y'}, axis=1)
daily_data_fb.head()
```

	ds	y
0	2011-12-31	4.424286
1	2012-01-01	4.677625
2	2012-01-02	4.991701
3	2012-01-03	5.175495
4	2012-01-04	5.125202

```
split_date = "2020-01-01"
train_split = daily_data_fb['ds'] <= split_date
test_split = daily_data_fb['ds'] > split_date
```

```
train_fb = daily_data_fb[train_split]
test_fb = daily_data_fb[test_split]
```

```
print("train data shape :", train_fb.shape)
print("test data shape :", test_fb.shape)
```

```
train data shape : (2924, 2)
test data shape : (455, 2)
```

```
model_prophet = Prophet()
for feature in exogenous_features:
    model_prophet.add_regressor(feature)
```

```
model_prophet.fit(df_train[["Timestamp", "Weighted_Price"] +
exogenous_features].rename(columns={"Timestamp": "ds",
"Weighted_Price": "y"}))

forecast = model_prophet.predict(df_valid[["Timestamp",
"Weighted_Price"] + exogenous_features].rename(columns={"Timestamp":
"ds"}))
```

```
forecast.head()
```

```
00:00:13 - cmdstanpy - INFO - Chain [1] start processing
```

```
00:00:15 - cmdstanpy - INFO - Chain [1] done processing
```

	ds	trend	yhat_lower	yhat_upper	trend_lower	trend_upper
0	2020-01-01	2702.877325	7072.301253	7700.919906	2702.877325	2702.877325
1	2020-01-02	2702.983911	7019.668240	7651.231829	2702.983911	2702.983911
2	2020-01-03	2703.090497	6874.553688	7501.087476	2703.090497	2703.090497
3	2020-01-04	2703.197083	6788.016202	7456.747590	2703.197083	2703.197083
4	2020-01-05	2703.303669	6881.145692	7508.624325	2703.303669	2703.303669

	Close_mean_lag3	Close_mean_lag3_lower	Close_mean_lag3_upper
0	3072.371173	3072.371173	3072.371173
1	3042.349141	3042.349141	3042.349141
2	2997.689186	2997.689186	2997.689186
3	2979.850021	2979.850021	2979.850021
4	3008.359587	3008.359587	3008.359587

	Close_mean_lag30	...	weekly	weekly_lower	weekly_upper
0	-0.260959	...	0.726059	0.726059	0.726059
1	-0.260709	...	-6.347854	-6.347854	-6.347854
2	-0.260368	...	-10.520293	-10.520293	-10.520293
3	-0.260110	...	7.412343	7.412343	7.412343
4	-0.260116	...	5.141867	5.141867	5.141867

	yearly_lower	yearly_upper	multiplicative_terms
0	-55.586668	-55.586668	0.0
1	-72.035715	-72.035715	0.0

2	-87.736168	-87.736168	0.0
3	-102.582359	-102.582359	0.0
4	-116.484028	-116.484028	0.0

	multiplicative_terms_lower	multiplicative_terms_upper	yhat
0	0.0	0.0	7390.766672
1	0.0	0.0	7330.359539
2	0.0	0.0	7197.632892
3	0.0	0.0	7133.479597
4	0.0	0.0	7189.611819

[5 rows x 121 columns]

```
df_valid["Forecast_Prophet"] = forecast.yhat.values
```

C:\Users\RISHAV RAJ SINGH\AppData\Local\Temp\ipykernel\_44116\3795914625.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

*# Plot Our Predictions*

```
fig1 = model_prophet.plot(forecast)
fig1.set_size_inches(15, 5)
plt.show()
```

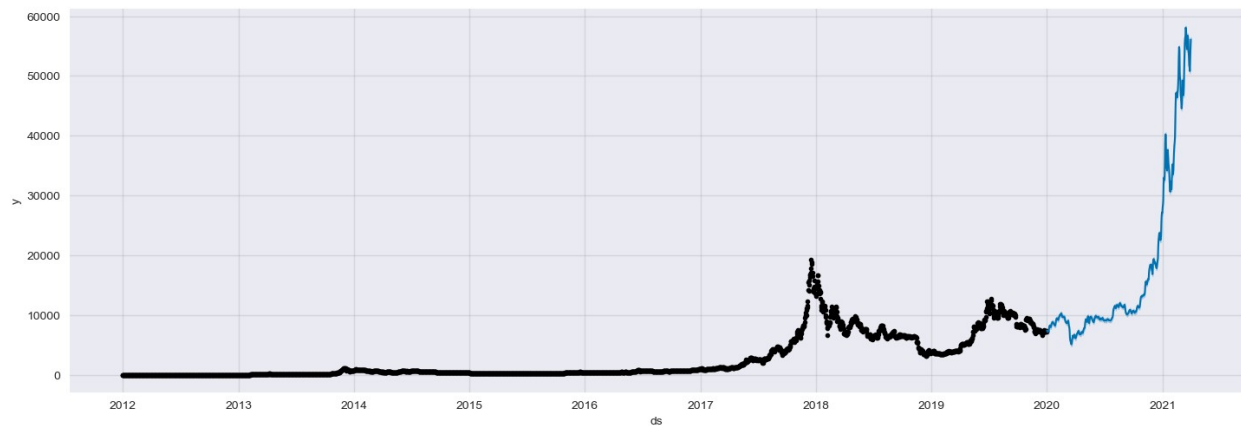
C:\Users\RISHAV RAJ SINGH\anaconda3\Lib\site-packages\prophet\plot.py:72: FutureWarning:

The behavior of `DatetimeProperties.to_pydatetime` is deprecated, in a future version this will return a Series containing python datetime objects instead of an ndarray. To retain the old behavior, call ``np.array`` on the result

C:\Users\RISHAV RAJ SINGH\anaconda3\Lib\site-packages\prophet\plot.py:73: FutureWarning:

The behavior of `DatetimeProperties.to_pydatetime` is deprecated, in a future version this will return a Series containing python datetime objects instead of an ndarray. To retain the old behavior, call


``np.array` on the result`



```
# Plotting changepoints
from prophet.plot import add_changepoints_to_plot

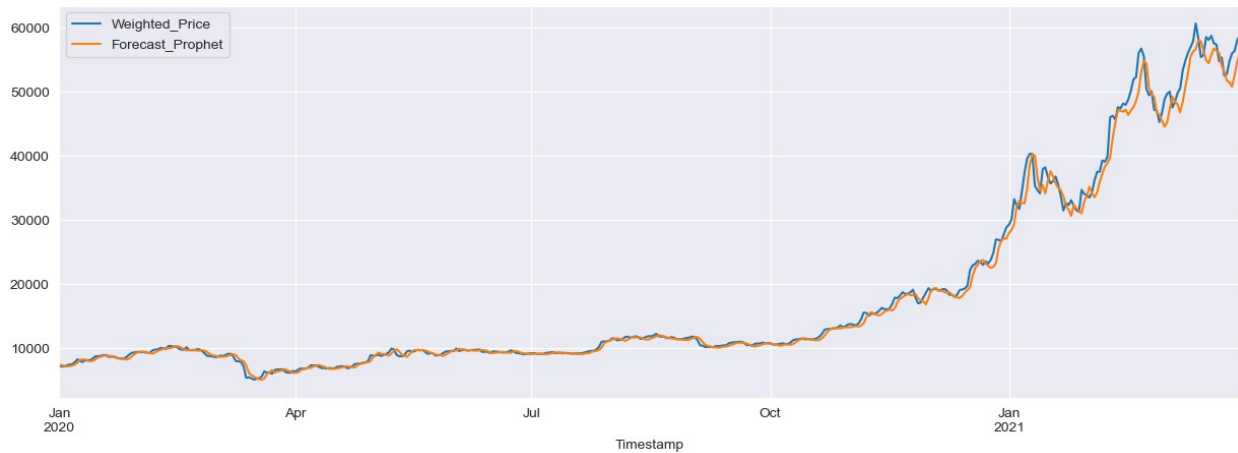
fig = model_prophet.plot(forecast)

a = add_changepoints_to_plot(fig.gca(), model_prophet, forecast)
```



```
df_valid[["Weighted_Price", "Forecast_Prophet"]].plot(figsize=(15, 5))
```

<Axes: xlabel='Timestamp'>



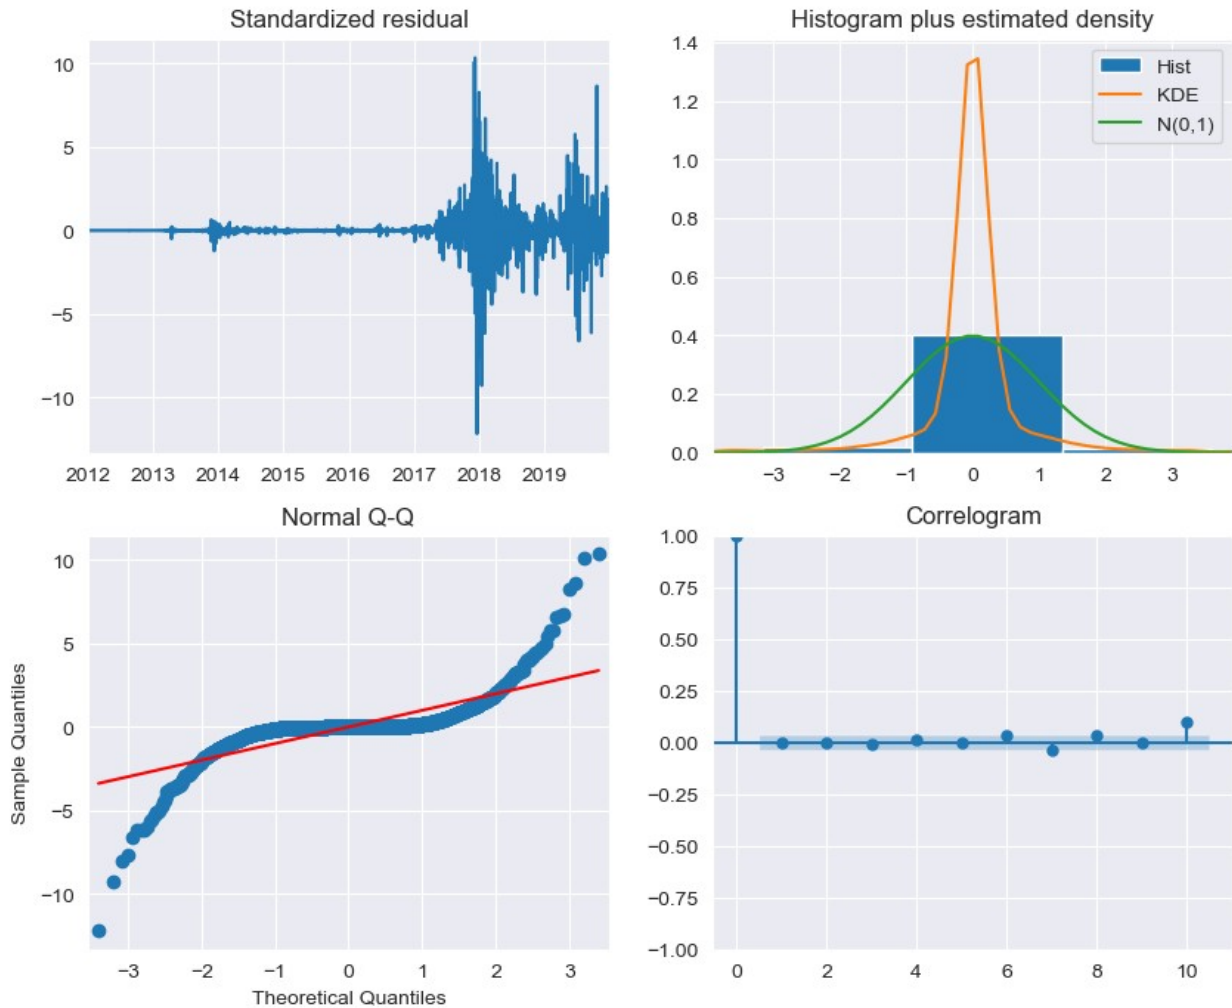
```
test_mae_prophet = mean_absolute_error(df_valid['Weighted_Price'],  
df_valid['Forecast_Prophet'])
```

```
test_rmse_prophet =  
np.sqrt(mean_squared_error(df_valid['Weighted_Price'],  
df_valid['Forecast_Prophet']))
```

```
print(f"Prophet's Test MAE : {test_mae_prophet}")  
print(f"Prophet's Test RMSE : {test_rmse_prophet}")
```

```
Prophet's Test MAE : 667.2904494828131  
Prophet's Test RMSE : 1250.55220793491
```

```
model.plot_diagnostics().set_size_inches(10, 8) # Adjust the size as  
needed  
plt.show()
```



#xgboost

```
!pip install xgboost
from sklearn import ensemble
from sklearn import metrics
from sklearn.model_selection import RandomizedSearchCV
import xgboost as xgb
from xgboost import plot_importance, plot_tree
from sklearn.metrics import mean_squared_error, mean_absolute_error
plt.style.use('fivethirtyeight')
```

```
from datetime import datetime
```

```
Requirement already satisfied: xgboost in c:\users\rishav raj singh\
anaconda3\lib\site-packages (2.0.3)
Requirement already satisfied: numpy in c:\users\rishav raj singh\
anaconda3\lib\site-packages (from xgboost) (1.26.4)
Requirement already satisfied: scipy in c:\users\rishav raj singh\
anaconda3\lib\site-packages (from xgboost) (1.11.4)
```





```

interaction_constraints=None,
                                learning_rate=... 0.8,
0.9, 1],
                                'colsample_bytree': [0.5, 0.6,
0.7, 0.8,
                                0.9, 1],
                                'gamma': [0.0, 0.1, 0.2, 0.3,
0.4, 0.5,
                                0.6, 0.7, 0.8, 0.9,
1],
                                'learning_rate': [0.05, 0.1,
0.15, 0.2,
                                0.25, 0.3],
                                'max_depth': [1, 3, 4, 5, 6,
7],
                                'min_child_weight': [3, 4, 5,
6, 7, 8,
                                9, 10,
11, 12, 13,
                                14],
                                'n_estimators': [100, 311,
522, 733,
                                944, 1155,
1366, 1577,
                                1788, 2000],
                                'subsample': [0.1, 0.2, 0.3,
0.4, 0.5,
                                0.6, 0.7, 0.8,
0.9, 1]},
                                verbose=3)

print(f"Model Best Score : {model.best_score}")
print(f"Model Best Parameters : {model.best_estimator_.get_params()}")

Model Best Score : -3.04475844490456
Model Best Parameters : {'objective': 'reg:squarederror',
'base_score': None, 'booster': None, 'callbacks': None,
'colsample_bylevel': 0.6, 'colsample_bynode': None,
'colsample_bytree': 0.5, 'device': None, 'early_stopping_rounds':
None, 'enable_categorical': False, 'eval_metric': None,
'feature_types': None, 'gamma': 0.3, 'grow_policy': None,
'importance_type': None, 'interaction_constraints': None,
'learning_rate': 0.15, 'max_bin': None, 'max_cat_threshold': None,
'max_cat_to_onehot': None, 'max_delta_step': None, 'max_depth': 6,
'max_leaves': None, 'min_child_weight': 3, 'missing': nan,
'monotone_constraints': None, 'multi_strategy': None, 'n_estimators':
1577, 'n_jobs': None, 'num_parallel_tree': None, 'random_state': None,
'reg_alpha': None, 'reg_lambda': None, 'sampling_method': None,

```

```
'scale_pos_weight': None, 'subsample': 0.9, 'tree_method': None,
'validate_parameters': None, 'verbosity': None}
```

```
model.best_estimator_
```

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=0.6, colsample_bynode=None,
              colsample_bytree=0.5,
              device=None, early_stopping_rounds=None,
              enable_categorical=False,
              eval_metric=None, feature_types=None, gamma=0.3,
              grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=0.15, max_bin=None, max_cat_threshold=None,
              max_cat_to_onehot=None, max_delta_step=None, max_depth=6,
              max_leaves=None, min_child_weight=3, missing=nan,
              monotone_constraints=None, multi_strategy=None,
              n_estimators=1577,
              n_jobs=None, num_parallel_tree=None,
              random_state=None, ...)
```

```
df_train['Predicted_Weighted_Price'] = model.predict(X_train)
```

```
df_train[['Weighted_Price', 'Predicted_Weighted_Price']].plot(figsize=(
15, 5))
```


```
C:\Users\RISHAV RAJ SINGH\AppData\Local\Temp\
ipykernel_44116\3944082182.py:1: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
<Axes: xlabel='Timestamp'>
```



```
df_valid['Forecast_XGBoost'] = model.predict(X_test)
```

```
overall_data = pd.concat([df_train, df_valid], sort=False)
```

```
C:\Users\RISHAV RAJ SINGH\AppData\Local\Temp\
ipykernel_44116\390820625.py:1: SettingWithCopyWarning:
```

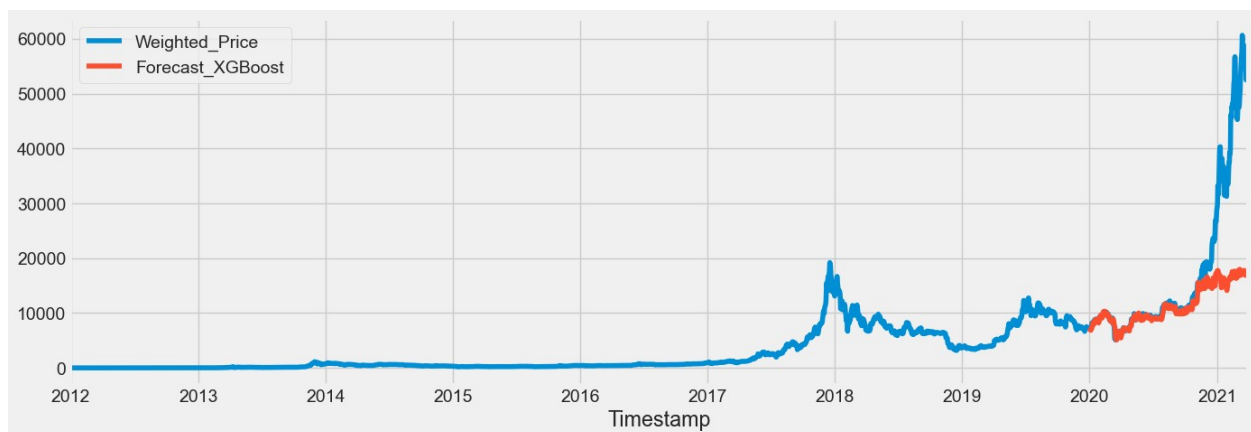
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)


```
overall_data[['Weighted_Price', 'Forecast_XGBoost']].plot(figsize=(15, 5))
```

```
<Axes: xlabel='Timestamp'>
```



```
df_valid[['Weighted_Price', 'Forecast_XGBoost']].plot(figsize=(15, 5))
```

<Axes: xlabel='Timestamp'>



```
from sklearn.metrics import mean_absolute_error,
mean_squared_error, r2_score

train_mae_xgboost = mean_absolute_error(df_train['Weighted_Price'],
df_train['Predicted_Weighted_Price'])
train_rmse_xgboost =
np.sqrt(mean_squared_error(df_train['Weighted_Price'],
df_train['Predicted_Weighted_Price']))
train_r_square_xgboost = r2_score(df_train['Weighted_Price'],
df_train['Predicted_Weighted_Price'])

print(f"train MAE : {train_mae_xgboost}")
print(f"train RMSE : {train_rmse_xgboost}")
print(f"train R2 : {train_r_square_xgboost}")

train MAE : 0.32486661540843664
train RMSE : 0.4374347732644935
train R2 : 0.999999985179853

test_mae_xgboost = mean_absolute_error(df_valid['Weighted_Price'],
df_valid['Forecast_XGBoost'])
test_rmse_xgboost =
np.sqrt(mean_squared_error(df_valid['Weighted_Price'],
df_valid['Forecast_XGBoost']))
test_r2_xgboost = r2_score(df_valid['Weighted_Price'],
df_valid['Forecast_XGBoost'])


print(f"test MAE XGB00ST : {test_mae_xgboost}")
print(f"test RMSE XGB00ST : {test_rmse_xgboost}")
print(f"test R2 XGB00ST : {test_r2_xgboost}")

test MAE XGB00ST : 6412.765432274534
test RMSE XGB00ST : 13319.975715568999
test R2 XGB00ST : 0.17193759741592907
```

```
#overall forecast of all model
```

```
df_valid[["Weighted_Price", "Forecast_ARIMAX", "Forecast_Prophet",  
"Forecast_XGBoost"]].plot(figsize=(15,7))
```

```
<Axes: xlabel='Timestamp'>
```



```
arimax_rmse = np.sqrt(mean_squared_error(df_valid['Weighted_Price'],  
df_valid['Forecast_ARIMAX']))
```

```
fbp_rmse = np.sqrt(mean_squared_error(df_valid['Weighted_Price'],  
df_valid['Forecast_Prophet']))
```

```
xgb_rmse = np.sqrt(mean_squared_error(df_valid['Weighted_Price'],  
df_valid['Forecast_XGBoost']))
```

```
# MAE
```

```
arimax_mae = mean_absolute_error(df_valid['Weighted_Price'],  
df_valid['Forecast_ARIMAX'])
```

```
fbp_mae = mean_absolute_error(df_valid['Weighted_Price'],  
df_valid['Forecast_Prophet'])
```

```
xgb_mae = mean_absolute_error(df_valid['Weighted_Price'],  
df_valid['Forecast_XGBoost'])
```

```
#mae,rmse of all model
```

```
print("ARIMAX RMSE :", arimax_rmse)
```

```
print("FB Prophet RMSE :", fbp_rmse)
```

```
print("XGBoost RMSE :", xgb_rmse)
```

```
print("\nARIMAX MAE :", arimax_mae)
```

```
print("FB Prophet MAE :", fbp_mae)
print("XGBoost MAE :", xgb_mae)

ARIMAX RMSE : 18052.527118342223
FB Prophet RMSE : 1250.55220793491
XGBoost RMSE : 13319.975715568999

ARIMAX MAE : 10692.217013361565
FB Prophet MAE : 667.2904494828131
XGBoost MAE : 6412.765432274534

print(f"Train RMSE Inverse: {LSTM_train_RMSE_inverse}")
print(f"Train MAE Inverse: {LSTM_train_MAE_inverse}")

print(f"Test RMSE Inverse: {LSTM_test_RMSE_inverse}")
print(f"Test MAE Inverse: {LSTM_test_MAE_inverse}")

Train RMSE Inverse: 879.906222025592
Train MAE Inverse: 774234.9595593504
Test RMSE Inverse: 6348.332969319255
Test MAE Inverse: 40301331.48934583
```